

```
In [35]: import tensorflow as tf
import tensorflow_decision_forests as tfdf
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Comment this if the data visualisations doesn't work on your side
%matplotlib inline
```

```
In [3]: train_file_path = "train.csv"
dataset_df = pd.read_csv(train_file_path)
print("Full train dataset shape is {}".format(dataset_df.shape))
```

Full train dataset shape is (1460, 81)

```
In [4]: dataset_df.head(10)
```

```
Out[4]:
```

	<b>Id</b>	<b>MSSubClass</b>	<b>MSZoning</b>	<b>LotFrontage</b>	<b>LotArea</b>	<b>Street</b>	<b>Alley</b>	<b>LotShape</b>	<b>LandCor</b>
<b>0</b>	1	60	RL	65.0	8450	Pave	NaN	Reg	
<b>1</b>	2	20	RL	80.0	9600	Pave	NaN	Reg	
<b>2</b>	3	60	RL	68.0	11250	Pave	NaN	IR1	
<b>3</b>	4	70	RL	60.0	9550	Pave	NaN	IR1	
<b>4</b>	5	60	RL	84.0	14260	Pave	NaN	IR1	
<b>5</b>	6	50	RL	85.0	14115	Pave	NaN	IR1	
<b>6</b>	7	20	RL	75.0	10084	Pave	NaN	Reg	
<b>7</b>	8	60	RL	NaN	10382	Pave	NaN	IR1	
<b>8</b>	9	50	RM	51.0	6120	Pave	NaN	Reg	
<b>9</b>	10	190	RL	50.0	7420	Pave	NaN	Reg	

10 rows × 81 columns


Warning: Total number of columns (81) exceeds max\_columns (20) limiting to first (20) columns.

```
In [5]: dataset_df = dataset_df.drop('Id', axis=1)
dataset_df.head(10)
```




Out[5]:


	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour
0	60	RL	65.0	8450	Pave	NaN	Reg	Lv
1	20	RL	80.0	9600	Pave	NaN	Reg	Lv
2	60	RL	68.0	11250	Pave	NaN	IR1	Lv
3	70	RL	60.0	9550	Pave	NaN	IR1	Lv
4	60	RL	84.0	14260	Pave	NaN	IR1	Lv
5	50	RL	85.0	14115	Pave	NaN	IR1	Lv
6	20	RL	75.0	10084	Pave	NaN	Reg	Lv
7	60	RL	NaN	10382	Pave	NaN	IR1	Lv
8	50	RM	51.0	6120	Pave	NaN	Reg	Lv
9	190	RL	50.0	7420	Pave	NaN	Reg	Lv

10 rows × 80 columns







## Distributions







Categorical distributions




2-d distributions




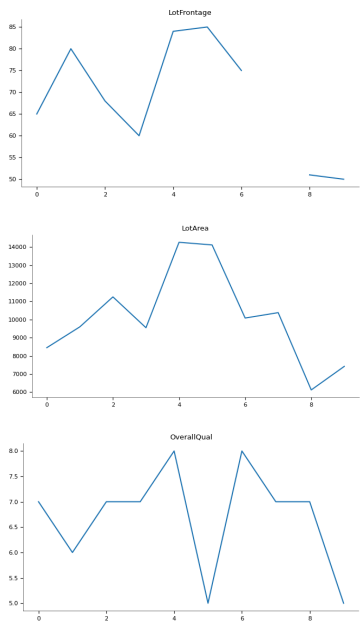


Time series





Values





2-d categorical distributions






## Faceted distributions


<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.




<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.




<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.



<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.



In [6]: dataset\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1460 entries, 0 to 1459
```

```
Data columns (total 80 columns):
```

#	Column	Non-Null Count	Dtype
0	MSSubClass	1460 non-null	int64
1	MSZoning	1460 non-null	object
2	LotFrontage	1201 non-null	float64
3	LotArea	1460 non-null	int64
4	Street	1460 non-null	object
5	Alley	91 non-null	object
6	LotShape	1460 non-null	object
7	LandContour	1460 non-null	object
8	Utilities	1460 non-null	object
9	LotConfig	1460 non-null	object
10	LandSlope	1460 non-null	object
11	Neighborhood	1460 non-null	object
12	Condition1	1460 non-null	object
13	Condition2	1460 non-null	object
14	BldgType	1460 non-null	object
15	HouseStyle	1460 non-null	object
16	OverallQual	1460 non-null	int64
17	OverallCond	1460 non-null	int64
18	YearBuilt	1460 non-null	int64
19	YearRemodAdd	1460 non-null	int64
20	RoofStyle	1460 non-null	object
21	RoofMatl	1460 non-null	object
22	Exterior1st	1460 non-null	object
23	Exterior2nd	1460 non-null	object
24	MasVnrType	588 non-null	object
25	MasVnrArea	1452 non-null	float64
26	ExterQual	1460 non-null	object
27	ExterCond	1460 non-null	object
28	Foundation	1460 non-null	object
29	BsmtQual	1423 non-null	object
30	BsmtCond	1423 non-null	object
31	BsmtExposure	1422 non-null	object
32	BsmtFinType1	1423 non-null	object
33	BsmtFinSF1	1460 non-null	int64
34	BsmtFinType2	1422 non-null	object
35	BsmtFinSF2	1460 non-null	int64
36	BsmtUnfSF	1460 non-null	int64
37	TotalBsmtSF	1460 non-null	int64
38	Heating	1460 non-null	object
39	HeatingQC	1460 non-null	object
40	CentralAir	1460 non-null	object
41	Electrical	1459 non-null	object
42	1stFlrSF	1460 non-null	int64
43	2ndFlrSF	1460 non-null	int64
44	LowQualFinSF	1460 non-null	int64
45	GrLivArea	1460 non-null	int64
46	BsmtFullBath	1460 non-null	int64
47	BsmtHalfBath	1460 non-null	int64
48	FullBath	1460 non-null	int64
49	HalfBath	1460 non-null	int64
50	BedroomAbvGr	1460 non-null	int64
51	KitchenAbvGr	1460 non-null	int64
52	KitchenQual	1460 non-null	object
53	TotRmsAbvGrd	1460 non-null	int64
54	Functional	1460 non-null	object

```

55 Fireplaces      1460 non-null  int64
56 FireplaceQu     770 non-null   object
57 GarageType      1379 non-null  object
58 GarageYrBlt     1379 non-null  float64
59 GarageFinish    1379 non-null  object
60 GarageCars      1460 non-null  int64
61 GarageArea      1460 non-null  int64
62 GarageQual      1379 non-null  object
63 GarageCond      1379 non-null  object
64 PavedDrive      1460 non-null  object
65 WoodDeckSF      1460 non-null  int64
66 OpenPorchSF     1460 non-null  int64
67 EnclosedPorch   1460 non-null  int64
68 3SsnPorch       1460 non-null  int64
69 ScreenPorch     1460 non-null  int64
70 PoolArea        1460 non-null  int64
71 PoolQC          7 non-null    object
72 Fence           281 non-null  object
73 MiscFeature     54 non-null    object
74 MiscVal         1460 non-null  int64
75 MoSold          1460 non-null  int64
76 YrSold          1460 non-null  int64
77 SaleType        1460 non-null  object
78 SaleCondition   1460 non-null  object
79 SalePrice       1460 non-null  int64
dtypes: float64(3), int64(34), object(43)
memory usage: 912.6+ KB

```

```

In [7]: print(dataset_df['SalePrice'].describe())
plt.figure(figsize=(9, 8))
sns.distplot(dataset_df['SalePrice'], color='g', bins=100, hist_kws={'alpha': 0.

```

```

count      1460.000000
mean      180921.195890
std        79442.502883
min        34900.000000
25%       129975.000000
50%       163000.000000
75%       214000.000000
max        755000.000000
Name: SalePrice, dtype: float64

```

<ipython-input-7-dc911a47893e>:3: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```

sns.distplot(dataset_df['SalePrice'], color='g', bins=100, hist_kws={'alpha':
0.4});

```





```
In [8]: list(set(dataset_df.dtypes.tolist()))
```


```
Out[8]: [dtype('float64'), dtype('int64'), dtype('O')]
```

```
In [9]: df_num = dataset_df.select_dtypes(include = ['float64', 'int64'])
df_num.head()
```

```
Out[9]:
```

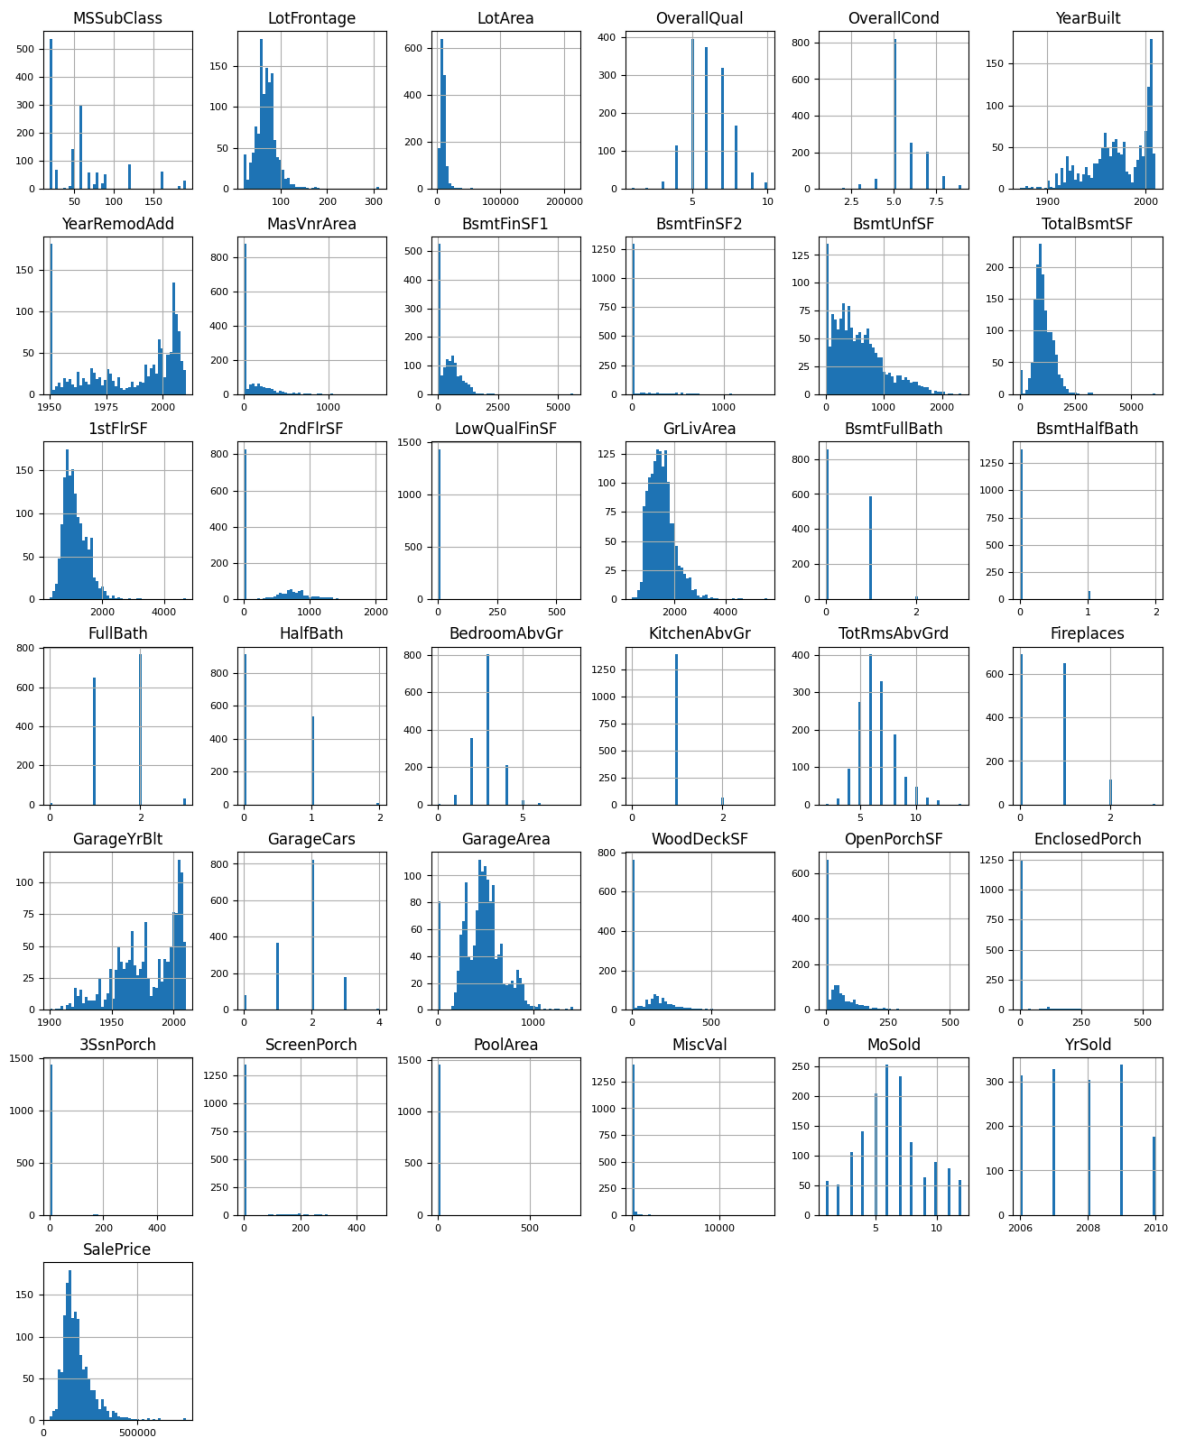
	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodA
0	60	65.0	8450	7	5	2003	2
1	20	80.0	9600	6	8	1976	1
2	60	68.0	11250	7	5	2001	2
3	70	60.0	9550	7	5	1915	1
4	60	84.0	14260	8	5	2000	2

5 rows × 37 columns



Warning: Total number of columns (37) exceeds max\_columns (20) limiting to first (20) columns.

```
In [10]: df_num.hist(figsize=(16, 20), bins=50, xlabelsize=8, ylabelsize=8);
```



```
In [11]: import numpy as np

def split_dataset(dataset, test_ratio=0.30):
    test_indices = np.random.rand(len(dataset)) < test_ratio
    return dataset[~test_indices], dataset[test_indices]

train_ds_pd, valid_ds_pd = split_dataset(dataset_df)
print("{} examples in training, {} examples in testing.".format(
    len(train_ds_pd), len(valid_ds_pd)))
```

1017 examples in training, 443 examples in testing.

```
In [12]: label = 'SalePrice'
train_ds = tfdf.keras.pd_dataframe_to_tf_dataset(train_ds_pd, label=label, task
valid_ds = tfdf.keras.pd_dataframe_to_tf_dataset(valid_ds_pd, label=label, task
```

```
In [13]: tfdf.keras.get_all_models()
```

```
Out[13]: [tensorflow_decision_forests.keras.RandomForestModel,  
          tensorflow_decision_forests.keras.GradientBoostedTreesModel,  
          tensorflow_decision_forests.keras.CartModel,  
          tensorflow_decision_forests.keras.DistributedGradientBoostedTreesModel]
```

```
In [14]: rf = tfdf.keras.RandomForestModel(task = tfdf.keras.Task.REGRESSION)  
         rf.compile(metrics=["mse"])
```

Use /tmp/tmp90d7frb1 as temporary training directory

```
In [15]: rf.fit(x=train_ds)
```


Reading training dataset...  
Training dataset read in 0:00:10.971011. Found 1017 examples.  
Training model...  
Model trained in 0:00:03.811701  
Compiling model...  
Model compiled.

```
Out[15]: <tf_keras.src.callbacks.History at 0x7aaca12398a0>
```

```
In [16]: tfdf.model_plotter.plot_model_in_colab(rf, tree_idx=0, max_depth=3)
```

```
Out[16]:
```

```
In [17]: import matplotlib.pyplot as plt  
         logs = rf.make_inspector().training_logs()  
         plt.plot([log.num_trees for log in logs], [log.evaluation.rmse for log in logs])  
         plt.xlabel("Number of trees")  
         plt.ylabel("RMSE (out-of-bag)")  
         plt.show()
```



```
In [18]: inspector = rf.make_inspector()
inspector.evaluation()
```

```
Out[18]: Evaluation(num_examples=1017, accuracy=None, loss=None, rmse=31358.87495781327,
ndcg=None, aucs=None, auuc=None, qini=None)
```

```
In [19]: evaluation = rf.evaluate(x=valid_ds,return_dict=True)

for name, value in evaluation.items():
    print(f"{name}: {value:.4f}")
```

```
1/1 [=====] - 9s 9s/step - loss: 0.0000e+00 - mse: 55074
5088.0000
loss: 0.0000
mse: 550745088.0000
```

```
In [20]: print(f"Available variable importances:")
for importance in inspector.variable_importances().keys():
    print("\t", importance)
```

```
Available variable importances:
    SUM_SCORE
    NUM_AS_ROOT
    INV_MEAN_MIN_DEPTH
    NUM_NODES
```

```
In [21]: inspector.variable_importances()["NUM_AS_ROOT"]
```

```
Out[21]: [("OverallQual" (1; #62), 96.0),
("ExterQual" (4; #22), 61.0),
("Neighborhood" (4; #59), 42.0),
("GarageCars" (1; #32), 38.0),
("GarageArea" (1; #31), 15.0),
("GrLivArea" (1; #38), 13.0),
("KitchenQual" (4; #44), 9.0),
("YearBuilt" (1; #76), 8.0),
("BsmtQual" (4; #14), 7.0),
("TotalBsmtSF" (1; #73), 4.0),
("1stFlrSF" (1; #0), 2.0),
("Foundation" (4; #28), 2.0),
("BsmtFinSF1" (1; #8), 1.0),
("GarageYrBlt" (1; #37), 1.0),
("YearRemodAdd" (1; #77), 1.0)]
```

```
In [22]: plt.figure(figsize=(12, 4))

# Mean decrease in AUC of the class 1 vs the others.
variable_importance_metric = "NUM_AS_ROOT"
variable_importances = inspector.variable_importances()[variable_importance_metr


# Extract the feature name and importance values.
#
# `variable_importances` is a list of <feature, importance> tuples.
feature_names = [vi[0].name for vi in variable_importances]
feature_importances = [vi[1] for vi in variable_importances]
# The feature are ordered in decreasing importance value.
feature_ranks = range(len(feature_names))

bar = plt.barh(feature_ranks, feature_importances, label=[str(x) for x in featur
plt.yticks(feature_ranks, feature_names)
```

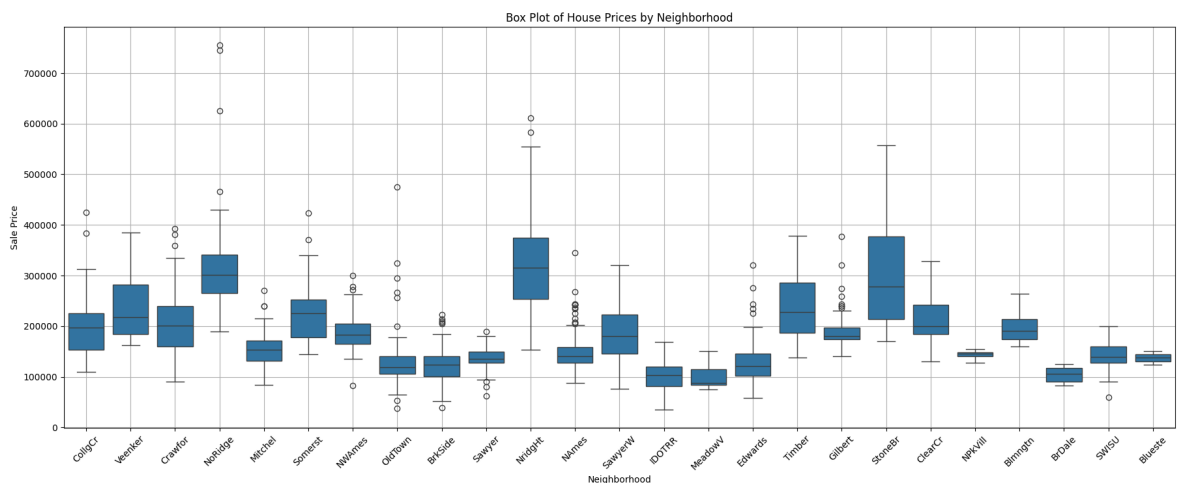
```
plt.gca().invert_yaxis()

# TODO: Replace with "plt.bar_label()" when available.
# Label each bar with values
for importance, patch in zip(feature_importances, bar.patches):
    plt.text(patch.get_x() + patch.get_width(), patch.get_y(), f"{importance:.4f}")

plt.xlabel(variable_importance_metric)
plt.title("NUM AS ROOT of the class 1 vs the others")
plt.tight_layout()
plt.show()
```



```
In [24]: plt.figure(figsize=(22, 8))
sns.boxplot(x='Neighborhood', y='SalePrice', data=dataset_df)
plt.xticks(rotation=45)
plt.title('Box Plot of House Prices by Neighborhood')
plt.xlabel('Neighborhood')
plt.ylabel('Sale Price')
plt.grid(True)
plt.show()
```





```
In [27]: numerical_cols = dataset_df.select_dtypes(include=['number']).columns
numerical_df = dataset_df[numerical_cols]


# Calculate the correlation matrix
correlation_matrix = numerical_df.corr()

# Plot the heatmap
plt.figure(figsize=(22, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()
```








## 2-d distributions



## Time series



## Values



```
In [34]: sample_submission_df = pd.read_csv('sample_submission.csv')
sample_submission_df['SalePrice'] = rf.predict(test_ds)
```

```
sample_submission_df.to_csv('submission.csv', index=False)
```

```
2/2 [=====] - 0s 28ms/step
```

```
In [ ]:
```